

# Online Algorithm and Online Optimization

---

Henry Vu

# Online Algorithm - An Example


- **Ski rental problem:**

At each day, player can decide to **rent** at \$1, or **buy** at \$B

If ski is bought, no need to rent (i.e. no future cost)

- **On day  $d$ , should you buy or rent?**

Need to know when the player stops skiing




# Ski Rental

- If stop date **D** is known beforehand
  - D**  $\geq$  **B**: buy on day 1
  - D**  $<$  **B**: rent every day

$$\text{Cost} = \min\{D, B\}$$

- But what if **D** is unknown?  
Need to design an **online** algorithm that minimizes cost.



# Competitive Online Algorithms

- Algorithm design where problem parameters  $\sigma$ 's are revealed sequentially
- **How to evaluate performance?** Compare how the online algorithm (ALG) compete vs its offline counterpart (OPT) that knows everything

**Competitive ratio:**  $\alpha = \frac{Cost(ALG)}{Cost(OPT)}$

- **Goal:** minimize  $\alpha$  in the worst case (also called worst-case competitive analysis)

# Competitive Ratios

- A **deterministic online algorithm** is called  $\alpha$ -competitive if:

$$\alpha = \max_{\sigma} \frac{ALG(\sigma)}{OPT(\sigma)}$$

- A **randomized online algorithm** (a set of deterministic algorithms over some probability distribution) is called  $\alpha$ -competitive if:

$$\alpha = \max_{\sigma} \frac{\mathbb{E}[ALG(\sigma)]}{OPT(\sigma)}$$

## Back to Ski Rental

- Problem parameter  $\sigma = D$
- **Offline algorithm:**  $\text{OPT}(D) = \min\{D, B\}$
- **Deterministic online algorithm:** decides to buy on some day **d** (renting up to day d-1)
  - if  $d \leq D$  i.e. buy ski before actually stop  
 $\text{ALG}(D) = d-1 + B$
  - if  $d > D$  i.e. rent until stop  
 $\text{ALG}(D) = D$

# Competitive Analysis

- In the worst case scenario, the player is forced to stop right after buying  
i.e.  $d = D$

$$\alpha = \max_D \frac{ALG(D)}{OPT(D)} = \max_D \frac{d - 1 + B}{\min(D, B)} = \frac{d - 1 + B}{\min(d, B)}$$

- The solution  $d^*$  that minimizes  $\alpha$  is  $d^* = B$   
 $ALG(D) = B - 1 + B = 2B - 1$

- The best possible **deterministic online algorithm** has almost **twice the cost!**

# Can we do better?

- Try **randomization**: at the start, choose algorithm  $ALG_d$  with probability  $p_d$  where  $ALG_d$ : rent for  $d-1$  days then buy on day  $d$

Assume  $\alpha$ -competitiveness:

$$\mathbb{E}[ALG(\sigma)] \leq \alpha \cdot OPT(\sigma)$$

$$\Rightarrow \sum_{i=1}^D (i - 1 + B) \cdot p_i + \sum_{i>D} D \cdot p_i \leq \alpha \cdot \min(D, B)$$



## An example: $B = \$4$

- We have the following LP

$$4p_1 + p_2 + p_3 + p_4 + p_5 + p_6 + \dots \leq 1.\alpha \quad \text{for } D = 1$$

$$4p_1 + 5p_2 + 2p_3 + 2p_4 + 2p_5 + 2p_6 + \dots \leq 2.\alpha \quad \text{for } D = 2$$

$$4p_1 + 5p_2 + 6p_3 + 3p_4 + 3p_5 + 3p_6 + \dots \leq 3.\alpha \quad \text{for } D = 3$$

$$4p_1 + 5p_2 + 6p_3 + 7p_4 + 4p_5 + 4p_6 + \dots \leq 4.\alpha \quad \text{for } D = 4$$

$$4p_1 + 5p_2 + 6p_3 + 7p_4 + 8p_5 + 5p_6 + \dots \leq 4.\alpha \quad \text{for } D = 5$$

$$4p_1 + 5p_2 + 6p_3 + 7p_4 + 8p_5 + 9p_6 + \dots \leq 4.\alpha \quad \text{for } D = 6$$

...

# An example: $B = \$4$

- Table of **competitive ratios**

	D = 1	D = 2	D = 3	D ≥ 4
ALG <sub>1</sub>	4/1	4/2	4/3	4/4
ALG <sub>2</sub>	1/1	5/2	5/3	5/4
ALG <sub>3</sub>	1/1	2/2	6/3	6/4
ALG <sub>4</sub>	1/1	2/2	3/3	7/4
ALG <sub>5</sub>	1/1	2/2	3/3	8/4

## An example: $B = \$4$

- End up with:
 
$$4p_1 + p_2 + p_3 + p_4 \leq 1.\alpha$$

$$4p_1 + 5p_2 + 2p_3 + 2p_4 \leq 2.\alpha$$

$$4p_1 + 5p_2 + 6p_3 + 3p_4 \leq 3.\alpha$$

$$4p_1 + 5p_2 + 6p_3 + 7p_4 \leq 4.\alpha$$
- $\alpha$  is minimized when all equalities hold
- In general, solving for  $p_d$  then  $\alpha$  based on the assumption that  $\sum p_d = 1$ , we obtain:

$$p_i = \left(\frac{B-1}{B}\right)^{B-i} \cdot \frac{\alpha}{B}$$

$$\alpha \simeq \frac{e}{e-1} = 1.582$$

**Better than the  
optimal deterministic algorithm  
for large B**

# Does randomization always improve competitiveness?

- How to establish a **lower bound** for the competitive ratio of randomized algorithms?
- How to prove that the competitive ratio of some randomized algorithm is the **best possible** for all randomized algorithms

Yao's Principle

# Yao's Principle

## Statement:

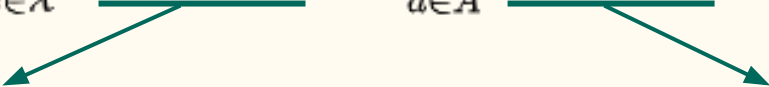
The expected cost of a **randomized algorithm** on the **worst-case input** is no better than the expected cost for a **worst-case probability distribution on the inputs** of the **deterministic algorithm** that performs best against that distribution.

# Yao's Principle: Proof

- The principle states that: Let  $\mathcal{X}$  and  $A$  be the set of all inputs and deterministic algorithms respectively.
- For any input  $x \in \mathcal{X}$  and algorithm  $a \in A$ , let  $c(a, x) \geq 0$  be the cost of running algorithm  $a$  on input  $x$ .
- Let  $X \in \mathcal{X}$  be a random input chosen according to some probability distribution  $p$ , and  $A_r$  be a random algorithm chosen according to some probability distribution  $q$ .

$$\max_{x \in \mathcal{X}} \mathbb{E}[c(A_r, x)] \geq \min_{a \in A} \mathbb{E}[c(a, X)]$$

## Yao's Principle: Proof (cont.)

$$\max_{x \in \mathcal{X}} \mathbb{E}[c(A_r, x)] \geq \min_{a \in A} \mathbb{E}[c(a, X)]$$


$\sum_a q_a c(a, x)$

$\sum_x p_x c(a, x)$

$$\begin{aligned}
 & \max_{x \in \mathcal{X}} \mathbb{E}[c(A_r, x)] \\
 = & \max_{x \in \mathcal{X}} \sum_a q_a c(a, x) \geq \sum_x p_x \sum_a q_a c(a, x) = \sum_a q_a \sum_x p_x c(a, x) = \sum_a q_a \mathbb{E}[c(a, X)] \\
 & \geq \min_{a \in A} \mathbb{E}[c(a, X)]
 \end{aligned}$$

## Another example: One-way trading

- When to trade your assets? Trade now or wait for better rates?
- **Example:** Consider a trader who wants to trade 1 CAD to USD. At time  $t$  ( $t = 1, 2, \dots, T$ ), rate  $v_t$  is revealed. Should the trader **wait** or **buy now**? If buy, how much to buy?

**Objective:** maximize the amount of USD





# Problem Formulation

- At each step  $t$ , trade  $x_t$  CAD for USD at rate  $v_t$
- Online Optimization Problem: input  $v_t$ 's are revealed sequentially

$$\begin{aligned} & \underset{x_t}{\text{minimize}} && \sum_{t=1}^T v_t x_t \\ & \text{subject to} && \sum_{t=1}^T x_t \leq 1, \\ & && x_t \geq 0, \quad \forall t \end{aligned}$$



# Competitive Analysis

- **Assumptions:**

Exchange rates (prices) are bounded:  $v_t \in [m, M]$  where  $0 < m \leq M$

$M$  and  $m$  are known to the online player

Can only sell 1 unit at each step

- The optimal deterministic algorithm is called **reservation price policy (RPP)**:  
Accept the first price greater or equal to some *reservation price*  $p^*$

- How to choose  $p^*$  that maximizes return?

## How to choose $p^*$ ?

- Consider 2 cases:

1. The search encounters some price  $\geq p^*$ . The worst case return is  $p^*$


$$\alpha = \max_{\sigma} \frac{OPT(\sigma)}{ALG(\sigma)} = \frac{M}{p^*}$$

2. The search doesn't encounter any price  $\geq p^*$ , thus, chooses the last price.  
The worst case return is  $m$

$$\alpha = \max_{\sigma} \frac{OPT(\sigma)}{ALG(\sigma)} = \frac{p^* - \epsilon}{m} \quad \text{for } \epsilon \rightarrow 0^+$$

## How to choose $p^*$ ? (cont.)

- The competitive ratio is the worst case of the 2:  $\alpha = \max \left\{ \frac{M}{p^*}, \frac{p^* - \epsilon}{m} \right\}$



So the optimal choice of  $p^*$  is  $p^* = \sqrt{Mm}$

# Online Primal Dual Approach

- Let  $p(\sigma)$  be the cost of the primal problem and  $d(\sigma)$  be the cost of its dual counterpart
- To achieve  $\alpha$ -competitiveness, we need  $\alpha = \max_{\sigma} \frac{p^*(\sigma)}{p_T(\sigma)}$  or  $p_T(\sigma) \geq \frac{1}{\alpha} p^*(\sigma)$
- By weak duality,  $p_t(\sigma) \geq \frac{1}{\alpha} d_t(\sigma) \geq \frac{1}{\alpha} d^*(\sigma) \geq \frac{1}{\alpha} p^*(\sigma)$

$$\sum_{t=1}^T (p_t(\sigma) - p_{t-1}(\sigma)) + p_0(\sigma)$$

$$\sum_{t=1}^T (d_t(\sigma) - d_{t-1}(\sigma)) + d_0(\sigma)$$

# Online Primal Dual Approach (cont.)

- So to ensure  $\alpha$ -competitiveness, we need to design a **trajectory** of feasible decision such that:  $\alpha \cdot \Delta_p \geq \Delta_d$
- **Challenges:**


How to design such trajectory?

How to minimize  $\alpha$ ?

How to satisfy the boundary conditions at each step  $t$  in the design?

# Online Knapsack Problem

- Classic (Offline) Knapsack Problem: Given a set of items, each with a weight and a value, choose some items such that their total weights is less than the total capacity of the knapsack while their total value is as large as possible.
- Online Knapsack Problem: The value and weight of each item is revealed **sequentially**. At each iteration  $t$ , determine whether or not to take the item.



# Online Knapsack: Assumptions

- **Assumption 1:** each item's weight is much smaller than the capacity of the knapsack, i.e.  $\exists \varepsilon$  close to 0 s.t  $w_t/B \leq \varepsilon$ .
- **Assumption 2:** the value-to-weight ratio of each item are bounded. i.e.  $\exists L, U \geq 0$  s.t.

$$L \leq \frac{v_t}{w_t} \leq U$$



# A simple deterministic algorithm

- Use **threshold**: If the item is “good enough”, i.e. its value-to-ratio is greater than some amount, pick that item.
- The algorithm becomes increasingly selective as knapsack capacity becomes scarcer.

For  $t = 1, 2, \dots, T$

If  $\frac{v_t}{w_t} \geq p_t$ , accept  $\longrightarrow x_t = 1$

Else if  $\frac{v_t}{w_t} < p_t$ , reject  $\longrightarrow x_t = 0$

$$p_t = \varphi(y_{t-1}) \quad \text{where } y_t = \sum_{t=1}^T w_t x_t$$

is an increasing function to reflect pickiness as more space is filled


# Intuitions about threshold price $p_t$

- At the beginning when the knapsack is empty ( $y \in [0, \beta]$ ), it could be that  $p_t \leq L \leq \frac{v_t}{w_t} \quad \forall t$ , i.e. we **accept all items** until the knapsack is filled to a certain point.
- If  $p_t$  increases too rapidly, and/or starts at some value  $\geq L$ , then

**OPT makes use of entire knapsack**

$$\alpha = \max_{\sigma} \frac{OPT(\sigma)}{ALG(\sigma)} \geq \frac{OPT(p_{min})}{ALG(p_{min})} = \frac{p_{min} \cdot 1}{p_{min} \cdot w_1} \rightarrow \infty$$


**ALG only able to accept first item**



## Intuitions about $p_t$ (cont.)

- We can henceforth assume  $p_t = L$  for  $y \in [0, \beta]$
- $p_t$  must be **non-decreasing** for  $y \in [\beta, 1]$  to reflect scarcity

How to design  $\beta$  and  $p_t$ ?



# Online Knapsack: LP Formulation

- Knapsack capacity is normalized to be 1. Item  $t$  ( $t = 1, 2, \dots, T$ ) has weight  $w_t$  and value  $v_t$ .

$$\begin{aligned} & \underset{x_t}{\text{maximize}} && \sum_{t=1}^T v_t x_t \\ & \text{subject to} && \sum_{t=1}^T w_t x_t \leq 1, \\ & && x_t \in \{0,1\} \quad \forall t \end{aligned}$$



$$\begin{aligned} & \underset{x_t}{\text{maximize}} && \sum_{t=1}^T v_t x_t \\ & \text{subject to} && \sum_{t=1}^T w_t x_t \leq 1, \\ & && 0 \leq x_t \leq 1 \quad \forall t \end{aligned}$$



$$\begin{aligned} & \underset{\lambda, \mu_t}{\text{minimize}} && \lambda + \sum_{t=1}^T \mu_t \\ & \text{subject to} && v_t - \lambda w_t \leq \mu_t \quad \forall t, \\ & && \lambda \geq 0, \\ & && \mu_t \geq 0 \quad \forall t \end{aligned}$$

**Primal Problem  $P(\sigma)$**

**Relaxed Primal  
Problem  $P_r(\sigma)$**

**Dual Problem  $D(\sigma)$**

# How to design? OPD Approach

From the complementary slackness condition:

$$\mu_t > 0 \Rightarrow x_t = 1 \quad \longrightarrow$$

$$\text{and } v_t - \lambda w_t - \mu_t < 0 \Rightarrow x_t = 0$$



$$\frac{v_t}{w_t} > \lambda \Leftrightarrow v_t - \lambda w_t \geq 0$$

From the inequality constraint of the dual problem

$$\mu_t \geq v_t - \lambda w_t \geq 0 \Rightarrow x_t = 1$$

$$\frac{v_t}{w_t} < \lambda \Leftrightarrow v_t - \lambda w_t < 0$$

From the inequality constraint of the dual problem

$$\mu_t \geq 0 \Rightarrow v_t - \lambda w_t - \mu_t < 0 \Rightarrow x_t = 0$$

## How to design? OPD Approach

- We design  $\lambda_t = p_t = \phi(y_t)$  where  $y_t$  captures resource utilization after step  $t$
- Since the primal optimal solution largely depends on the dual optimal solution, we rely on  $\lambda_{t-1}$  to design the online primal decision  $x_t$
- $$x_t = \begin{cases} 1, & \frac{v_t}{w_t} \geq \lambda_{t-1} = \phi(y_{t-1}) \\ 0, & \frac{v_t}{w_t} < \lambda_{t-1} = \phi(y_{t-1}) \end{cases}$$
- In addition, we design  $\mu_t = \max\{0, v_t - \lambda_{t-1}w_t\}$

## How to design? OPD Approach

- Now what we have left is to ensure **incremental inequality** holds  $\alpha.\Delta_p \geq \Delta_d$
- To achieve this, we need to design  $\phi$  such that

$$\phi(y_t) \leq \alpha\phi(y) \quad \text{for } y \in [0, 1]$$

# Design Principles

- There is a critical threshold  $\beta$  such that:  
 $\alpha \geq 1/\beta$   
 $\varphi(y) = (\text{or } \leq) L$  for  $y \in [0, \beta]$
- For the increasing segment  $y \in [\beta, 1]$ :  
 $\varphi'(y) \leq \alpha \varphi(y)$
- $\varphi(1) \geq U$



# Optimal Threshold Design

$$\phi(y) = \left(\frac{Ue}{L}\right)^y \left(\frac{L}{E}\right)$$


$$\beta = \frac{1}{1 + \ln\left(\frac{U}{L}\right)}$$

$$\alpha = \frac{1}{\beta} = 1 + \ln\left(\frac{U}{L}\right)$$

Threshold ALG  
achieves constant  
competitive ratio

**Best possible among all online  
algorithm!**

**Proof:** Yao's Principle



# Online Convex Optimization

- **Online Convex Optimization (OCO)**: rich theoretical area with a wide range of important applications.

In an OCO problem, an agent interacts with the environment in sequence.

At each time  $t$  ( $t = 1, 2, \dots, T$ ):

- i, The agent chooses action  $x_t$
- ii, The environment reveals cost function  $c_t$
- iii, The agent experiences cost  $c_t(x_t)$

**Objective:** Minimizes the *cumulative cost* over  $T$  ( $T$  is large) rounds

- Example: The well-known  $k$ -experts problem

# K-Experts Problem

- At each time  $t$ , the agent has to choose action A or B (e.g. to buy or sell stocks). There are  $N$  experts that offer advice. After a decision is made, the correct action is revealed and the agent accumulates some loss (let's say 0 if the agent chose the correct decision and 1 otherwise)
- **Observation:** Let  $L \geq T/2$  be the number of mistakes the best expert makes. Then there does not exist a deterministic algorithm that can guarantee less than  $2L$  mistakes.

# K-Experts Problem with Deterministic Algorithm

- **Proof:** Assume that there are 2 experts and one always chooses A and the other always chooses B. Since our algorithm is deterministic, the adversary can always choose a sequence such that our total mistakes at the end is  $T$ . In this settings, the best expert makes no more than  $T/2$  mistakes  
 $\Rightarrow$  A deterministic algorithm makes at least  $2L$  mistakes
- This observation motivates the design of a **randomized algorithm** for this problem.

# The Weighted Majority Algorithm

- **Intuition:** Each expert is initially assigned with some weight. At iteration  $t$ , if the total weights of all experts that choose A ( $W_t(A)$ ) is  $\geq$  the total weights of all experts that choose B ( $W_t(B)$ ), the agent chooses A (or B if  $W_t(B) \geq W_t(A)$ ). If an expert is wrong at iteration  $t$ , they have less influence on the agent's decision in the future ( $t+1, t+2, \dots, T$ )

- Specifically, for expert  $i$ ,  $w_{t+1}(i) = \begin{cases} w_t(i), & \text{if } i \text{ was right} \\ (1 - \epsilon)w_t(i), & \text{if } i \text{ was wrong} \end{cases}$

- $x_t = \begin{cases} A, & \text{if } W_t(A) \geq W_t(B) \\ B, & \text{otherwise} \end{cases}$  where  $W_t(A) = \sum_{i \text{ that chose A}} w_t(i)$   
and  $W_t(B) = \sum_{i \text{ that chose B}} w_t(i)$

# Performance Bounds

- **Theorem 1:** Supposed  $\varepsilon \in [0, \frac{1}{2}]$  and the best expert makes  $L$  mistakes and the algorithm makes  $M$  mistakes. Then:

1. There is an efficient **deterministic algorithm** that can guarantee

$$M \leq 2(1+\varepsilon)L + 2\log N/\varepsilon$$

2. There is an efficient **randomized algorithm** that can guarantee

$$E[M] \leq (1+\varepsilon)L + \log N/\varepsilon$$

This randomized algorithm is called **Randomized Weighted Majority** (RWM). In RWM, the agent chooses expert  $i$  at time  $t$  with probability

$$p_t(i) = \frac{w_t(i)}{\sum_{j=1}^N w_t(j)}$$

# Online Convex Optimization

- General framework:

**Input:** set of possible decisions (convex set  $C$ )

**For**  $t = 1, 2, \dots, T$ :

- . Make decision  $x_t \in C$
- . Loss is revealed after decision is made  $f_t(x_t)$

**Objective:** Minimize the cumulative loss:

$$\min_{x_t \in C} \sum_{t=1}^T f_t(x_t)$$

## Two Communities, Two Metrics

- The goal of the OCO community is most typically to develop algorithms that perform well with respect to **regret**: the difference between the cost of the online algorithm and the cost of the offline optimal fixed (static solution).

$$Regret = \sum_{t=1}^T f_t(x_t) - \min_{u \in C} \sum_{t=1}^T f_t(u)$$

- The online algorithm community focuses on **competitive ratio**, which is the worst-case ratio between the cost of the online algorithm and the cost of the optimal offline solution (dynamic solution).

$$\alpha = \max_{\sigma} \frac{ALG(\sigma)}{OPT(\sigma)}$$



## Two Communities, Two Metrics

- The variation of OCO that the Online Algorithm community focuses on differs in that the cost  $c_t(x_t)$  is revealed to the agent at iteration  $t$  before the agent makes a decision. This is the main difference between the two communities  
optimal dynamic solution (CR) vs optimal static solution (Regret)
- It is desirable for an algorithm to learn both dynamic and static concepts